

Khyati Parihar

San Diego, CA | (858) 214-0846 | khyatiparihar@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Senior analyst with 6+ years of experience turning ML/AI product data into insights that improve adoption, model performance, retention, and business outcomes. Skilled in stakeholder management, ROI analysis, and scalable data pipeline design with experience spanning AWS cloud infrastructure, API integration, and end-to-end product lifecycle management

EDUCATION

University of California San Diego, La Jolla, California **06/2026**

M.S. Business Analytics (Rady Incoming Scholarship Award Recipient)

Relevant Coursework - Deep Learning & Neural Networks, NLP, Recommender Systems, Large-Scale ML Systems, Causal Inference, Cloud Computing

University of Delhi, New Delhi, India **05/2019**

Bachelor of Arts in Economics (Honors)

SKILLS

Technical: Python (Pandas, NumPy, scikit-learn, PyTorch, XGBoost), SQL, R, Git, Excel

Visualization & BI: Matplotlib, Tableau, Power BI, AWS QuickSight, Google Analytics

Analytics: A/B Testing, KPI Reporting, Cohort Analysis, Segmentation, Root-Cause Analysis, Statistical Analysis, EDA

ML & AI: Conversational AI Development, RAG (Retrieval-Augmented Generation), NLP, LLM Fine-Tuning, Classification, Regression, Feature Engineering, Hyperparameter Tuning

Automation: Prompt Engineering, Workflow Orchestration (n8n, Zapier), Claude API, HubSpot Breeze AI

Leadership & Strategy: Stakeholder Management, Cross-functional Collaboration, Client Onboarding, Technical Consulting

WORK EXPERIENCE

Senior ML Analyst, 80/20 Endurance, San Diego, CA **03/2026 – 06/2026**

- Led cross-functional capstone team as client-facing ML Consultant, partnering with C-suite stakeholders to diagnose operational bottlenecks, define automation requirements, and translate business requirements into agentic workflow solutions
- Diagnosed a broken customer operations process handling 23,000+ monthly tasks manually; designed and deployed end-to-end agentic AI system (HubSpot + Zapier + Claude API) that automated 95% of requests, reducing turnaround time from 12 hours to 2 minutes and achieving 99% task accuracy
- Architected and maintained scalable cross-system data pipelines and CI/CD workflows spanning APIs, CRM platforms, and communication tools, reducing operational turnaround time and improving process visibility across teams
- Created Python-based pre-production synthetic data testing framework, simulating thousands of edge-case scenarios across multiple LLMs to evaluate decision quality and prevent production failures

AI Consultant, Netomi **09/2024 – 06/2025**

- Built and deployed conversational AI systems for 40+ enterprise clients; fine-tuned Llama 2/3/3.1 and GPT-4 on millions of customer interactions, implemented RAG-based knowledge retrieval (AnswerNet) leveraging vector database architecture, and integrated into production via APIs
- Onboarded enterprise clients end-to-end by transforming raw conversation data into production-ready intent classification models, including model development, validation testing, and cross-functional deployment coordination
- Ran model comparison experiments across base models and fine-tuned variants, tracked classification accuracy, intent coverage, deflection rate, and CSAT as core KPIs, to decide which model version shipped for each client
- Established MLOps best practices by automating Python validation pipelines for pre-production QA, reducing client onboarding from 4+ weeks to under 2 weeks while maintaining 90%+ response accuracy
- Built the reporting layer for weekly bot performance, giving clients, CSMs, product owners, and internal teams a shared view of key metrics and improvement areas

ML Data Analyst, Netomi **07/2021 – 08/2024**

- Partnered with product, operations, and client success teams to support end-to-end analytics for intent classification and NLU deployments, spanning data preparation, KPI definition, evaluation, launch readiness, and post-launch monitoring
- Built executive-facing SQL and Tableau dashboards tracking deflection rate, CSAT, feature adoption, and model performance across 40+ deployments; used cohort and segmentation analysis to drive product improvements increasing client retention by 50%
- Performed root-cause analysis on production metrics using SQL and Python, translating performance trends into recommendations that improved workflows, model behavior, and client outcomes
- Optimized LLM model efficiency through intent decomposition and training data reduction, splitting overloaded intent clusters into targeted sub-models, improving response accuracy while reducing per-model resource consumption

- Drove analytics-led improvements across 40+ deployments; increased CSAT from 80% to 95%+, and translated performance data into executive-level recommendations that shaped roadmap and client investment decisions

Associate Analyst, Coforge Ltd.

11/2019 – 06/2021

- Developed SQL-driven Tableau and AWS QuickSight dashboards tracking weekly security KPIs across multiple business units; translated complex data into clear stakeholder narratives, accelerating decision-making by 30%
- Created Excel-based vulnerability tracking models with pivot tables and trend analysis to monitor PCI compliance release scans across TrustWave tools (Fusion, AppScanner, TrustKeeper), ensuring 100% regulatory compliance for secure releases
- Queried and analyzed vulnerability data using SQL to identify open threat patterns across OpCos; visualized findings in Tableau to prioritize remediation, reducing security exposure by 40% and accelerating remediation timelines

TECHNICAL PROJECTS

Intuit QuickBooks Upsell Targeting - ML Classification & Profit Optimization, Academic project

10/2025

- Engineered predictive features and trained Logistic Regression, Random Forest, XGBoost, & MLP ensemble models on 75K business records; applied profit-threshold targeting to scale predictions to an 800K+ population – ranked #2 among 60+ competing teams

Sydney Airport - Applied AI Hackathon, Professional project

11/2024

- Designed a GenAI solution for Sydney Airport, combining LLM prompting with real-time weather and flight delay data to improve traffic flow, reduce predicted passenger congestion by 30% – ranked #1 among 6 competing teams

United Airlines - Data Science Hackathon, Professional project

08/2024

- Built United Airlines' first LLM chatbot, optimizing classifier performance and response quality across deployment topics; used SQL and prompt-based extraction from airline web content to achieve 98% accuracy